

UNIVERSIDADE FEDERAL DE MINAS GERAIS
Instituto de Ciências Exatas
Programa de Pós-Graduação em Ciência da Computação

Tomaz Gomes Mascarenhas

Demonstrando teoremas em Lean por meio da reconstrução de provas em
SMT

Belo Horizonte
2023

Tomaz Gomes Mascarenhas

**Demonstrando teoremas em Lean por meio da reconstrução de provas em
SMT**

Versão Final

Dissertação apresentada ao Programa de Pós-Graduação em
Ciência da Computação da Universidade Federal de Minas
Gerais, como requisito parcial à obtenção do título de Mestre
em Ciência da Computação.

Orientador: Haniel Barbosa

Belo Horizonte
2023

Tomaz Gomes Mascarenhas

Proving Lean theorems via reconstructed SMT proofs

Final Version

Thesis presented to the Graduate Program in Computer Science of the Federal University of Minas Gerais in partial fulfillment of the requirements for the degree of Master in Computer Science.

Advisor: Haniel Barbosa

Belo Horizonte
2023

[Ficha Catalográfica em formato PDF]

A ficha catalográfica será fornecida pela biblioteca. Ela deve estar em formato PDF e deve ser passada como argumento do comando `ppgccufmg` no arquivo principal `.tex`, conforme o exemplo abaixo:

```
\ppgccufmg{  
    ...  
    fichacatalografica={ficha.pdf}  
}
```

[Folha de Aprovação em formato PDF]

A folha de aprovação deve estar em formato PDF e deve ser passada como argumento do comando `ppgccufmg` no arquivo principal `.tex`, conforme o exemplo abaixo:

```
\ppgccufmg{  
    ...  
    folhadeaprovacao={folha.pdf}  
}
```

Dedicuum cest laborae a quelquis personatum que ajudorat a facirelo.

Acknowledgments

Lorem ipsum dolor sit amet, consectetur adipisicing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

Sed ut perspiciatis unde omnis iste natus error sit voluptatem accusantium doloremque laudantium, totam rem aperiam, eaque ipsa quae ab illo inventore veritatis et quasi architecto beatae vitae dicta sunt explicabo. Nemo enim ipsam voluptatem quia voluptas sit aspernatur aut odit aut fugit, sed quia consequuntur magni dolores eos qui ratione voluptatem sequi nesciunt. Neque porro quisquam est, qui dolorem ipsum quia dolor sit amet, consectetur, adipisci velit, sed quia non numquam eius modi tempora incidunt ut labore et dolore magnam aliquam quaerat voluptatem. Ut enim ad minima veniam, quis nostrum exercitationem ullam corporis suscipit laboriosam, nisi ut aliquid ex ea commodi consequatur? Quis autem vel eum iure reprehenderit qui in ea voluptate velit esse quam nihil molestiae consequatur, vel illum qui dolorem eum fugiat quo voluptas nulla pariatur?

At vero eos et accusamus et iusto odio dignissimos ducimus qui blanditiis praesentium voluptatum deleniti atque corrupti quos dolores et quas molestias excepturi sint occaecati cupiditate non provident, similique sunt in culpa qui officia deserunt mollitia animi, id est laborum et dolorum fuga. Et harum quidem rerum facilis est et expedita distinctio. Nam libero tempore, cum soluta nobis est eligendi optio cumque nihil impedit quo minus id quod maxime placeat facere possimus, omnis voluptas assumenda est, omnis dolor repellendus. Temporibus autem quibusdam et aut officiis debitis aut rerum necessitatibus saepe eveniet ut et voluptates repudiandae sint et molestiae non recusandae. Itaque earum rerum hic tenetur a sapiente delectus, ut aut reiciendis voluptatibus maiores alias consequatur aut perferendis doloribus asperiores repellat.

“Só quem sonha acordado vê o sol nascer.”
(Unknown)

Resumo

Apesar de sua expressividade e robustez, assistentes de demonstração podem ser proibitivamente custosos para serem usados em formalizações de grande escala, dada a dificuldade de produzir as demonstrações interativamente. Atribuir a responsabilidade de demonstrar algumas das proposições a demonstradores automáticos de teoremas, como solucionadores de satisfatibilidade modulo teorias (SMT), é um jeito reconhecido de melhorar a usabilidade de assistentes de demonstração. Essa dissertação descreve uma nova integração entre o assistente de demonstração Lean 4 e o solucionador SMT cvc5.

Dada uma codificação de um teorema declarado em Lean como um problema de SMT e uma demonstração provida pelo cvc5 para o problema codificado, nós mostramos como traduzir essa demonstração para uma que certifique o teorema original em Lean. Para isso é necessário demonstrar a corretude, em Lean, dos passos lógicos tomados pelo solucionador. Desse modo, o verificador de demonstrações de Lean aceitará a demonstração em SMT do teorema original, caso o processo seja bem sucedido.

Essas técnicas são integradas no projeto Lean-SMT, que tem como objetivo criar uma tática em Lean que implemente o processo completo, isto é, a partir de um teorema em Lean, traduzi-lo para um problema formulado na linguagem do solucionador, invocar um solucionador para tentar resolvê-lo e produzir uma demonstração, e, caso ele seja bem-sucedido, traduzi-la para certificar o teorema original em Lean (o que é feito pelas técnicas apresentadas aqui).

Palavras-chave: Verificação Formal, Lean, SMT

Abstract

Despite their expressivity and robustness, interactive theorem provers (ITPs) can be prohibitively costly to use in large-scale formalizations due to the burden of interactively proving goals. Discharging some of these goals via automatic theorem provers, such as satisfiability modulo theories (SMT) solvers, is a known way of improving the usability of ITPs. This thesis describes a novel integration between the ITP Lean 4 and the SMT solver `cvc5`.

Given the encoding of some Lean goal as an SMT problem and a proof from `cvc5` of the encoded problem, we show how to lift this proof into a proof of the original goal. This requires proving the correctness, inside Lean, of the steps taken by the solver. Thus Lean’s proof checker will accept the SMT proof as a proof of the original goal, in case this process is successful.

This set of techniques is part of the joint project Lean-SMT, which aims to create a tactic in Lean that implements the whole pipeline, that is, from a goal in Lean, translate it into a query in the solver’s language, try to prove it using a solver and produce a proof and, in case it is successful, lift the proof produced, closing the original goal in Lean (which is done by the techniques presented here).

Keywords: Formal Verification, Lean, SMT

Contents

1	Introduction	11
1.1	Context	11
1.2	Contributions	13
1.3	Related Work	14
1.4	Organization of this document	16
2	Formal Preliminaries	17
2.1	Satisfiability Modulo Theories	17
2.2	Lean’s Type Theory	17
2.3	Lean’s Framework for Metaprogramming	17
3	Certifying Reconstruction of SMT Proofs in Lean	18
3.1	Certified vs Certifying	18
3.2	Classical vs Intuitionist (?)	18
3.3	Tactics	18
3.4	The Complete Architecture	18
3.5	Skipping the Parser	18
4	Evaluation	19
5	Future Work	20
	Bibliography	21

Chapter 1

Introduction

1.1 Context

A mechanized proof is a proof, written in some language recognized by a computer, that had its validity checked by a trusted verifier. One of the main applications of these artifacts are formalizing mathematical theories. Indeed, there are well-known examples of successful formalizations. One of them is the mechanization of the proof of a theorem regarding Perfectoid Spaces[10], done by the fields-medalist mathematician Peter Scholze, together with the community of a system called Lean[11]. Scholze proved the theorem using pen and paper, but was unsure of the result due to its complexity. Once he translated the theorem and the proof to the language of Lean, the system could point out some mistakes he made, and, after fixing them, he could be sure of the correctness of the proof.

Another application of mechanized proofs is verifying the correctness of mission-critical software. Given a specification of the behavior of some program, the program is said to be correct if it respects the specification for any input it is given. For instance, one could specify that a sorting routine must always produce the sorted permutation of its input list. In this case, a given sorting routine is said to be correct if it indeed produces the desired permutation, regardless of which list it receives. There are a variety of techniques to obtain correctness evidence for a software. The most common one is the development of tests. Besides being easy to write an efficient set of tests, there are many types of bugs that can be discovered with its execution. In fact, this approach is enough for a large amount of problems that are solved by software engineering. However, tests can't guarantee that a program doesn't have flaws, since the number of valid inputs is almost always exceedingly large, or infinite. This kind of guarantee is extremely important for mission-critical software, that is, systems that have critical responsibilities, such as the control of airplanes or medical equipment. In this context, one promising alternative is to use a mechanized proof of the correctness of the software as an evidence for its safety.

The process of generating mechanized proofs can be divided into two categories: interactive and automatic.

Interactive theorem provers (ITPs) are mainly represented by proof assistants, in which, after defining a theorem, the user attempts to manually write a proof for it, relying on the tool to organize the set of hypothesis and how the goal changed step-wisely through the proof, as well as to ensure the correctness of each step according to a trusted kernel. In order to keep the kernel simple and small (and, therefore, easy to be trusted), it's implementation usually just straightforwardly checks the logic rules from the logic implemented by the ITP. Because of this, each step must be explicitly stated by the user, making the tool costly to be used.

Automatic theorem provers (ATPs), on the other hand, only require the user to define a conjecture, proceeding automatically to determine whether there exists a proof for it, or possibly providing a counter-example if it can find one. Although they are easier to use, ATPs require a large codebase to implement all the algorithms necessary to execute the search for a proof, making them more susceptible to errors and harder to be trusted. One possible way to overcome this trust issue is to produce a mechanized proof verifying the correctness of the ATP, however, besides being a very complex task, once the proof is done the development of the ATP becomes freezed, otherwise it would have to be verified again.

Another approach to increase the confidence in ATPs is to have them provide a proof to support their results, so that it can be independently verified whether it indeed proves the theorem in question. This has the downside of creating a need for allocating resources to verify the proof of every single theorem that is proved. On the other hand, as long as the proof format doesn't change, the implementation of the solver can be modified without requiring a modification in the checkers. Also, it is important to consider that it is often simpler to verify proofs than to verify the tool itself.

Another important advantage of the second approach is that it allows the ITPs to leverage the automatic proving performed by the ATPs by using the proofs they produce, since the requirement for accepting a proof, i.e. that each step is correct to its internal logic, can be applied to the ATP proof. By connecting these systems, it would be possible for the user of the ITP to focus on more complex steps of the proof, such as defining an induction hypothesis, while delegating the burden of other long and straightforward steps to the ATP. Indeed, this connection is so important that there are projects like *Hammering Towards QED* [6] that outline all the efforts that were already made in order to integrate interactive and automatic theorem provers. In this paper, the authors describe in detail each component that a system that creates a connection between ATPs and ITPs has to implement, as well as the main issues that they have to solve, based on existing programs that were successful in this task. Besides that, they show their potential through several large benchmarks.

1.2 Contributions

Given this context, we present a set of tools that would be an essential part of the integration between the ITP Lean 4 [11] and the SMT solver `cvc5` [1]. Specifically, we aim to build a system that takes proofs of the unsatisfiability of SMT queries produced by `cvc5` and reconstructs them in Lean. The main motivation of this project is that despite the fact that Lean is emerging as a promising programming language and proof assistant and being widely used by mathematicians in large-scale formalizations [13, 10], there is currently no way to interact with SMT solvers from it, even though these systems have been central in previous developments of proof automation in ITPs, as we will show in Sections 1.3.2 and 1.3.3. The contribution of the present work would enable a faster development of this kind of project using Lean.

We use the `cvc5` solver because it already has a module for exporting proofs as Lean scripts [2], using a representation of the SMT terms¹ as an inductive type in Lean. However, these proofs are not fully verified by Lean’s checker. Instead, a set of axioms are declared in Lean, representing all the logical rules that `cvc5` uses to prove theorems, and the ITP only checks whether the rules were applied correctly and whether the end result of applying all the rules in the proof is, indeed, the required one. Our main contributions are to eliminate the need of increasing the trusted base by introducing those axioms and to make the proofs operate over native Lean terms, as opposed to terms of the inductive type that represents SMT terms.

Note that the set of tools we are proposing does not implement the full integration between Lean and `cvc5`. For instance, we do not implement a module for translating Lean goals into an equivalent SMT problem. However, our project is being used as part of the joint project Lean-SMT², that aims to implement a tactic in Lean that would perform the complete process, that is, starting from a Lean goal, translating it to a SMT query, invoking a solver to try to prove it and lifting the proof produced (in case it is found) to Lean’s language, so that it can be used as a proof for the original goal.

¹For more details about the SMT term language, see SMT-LIB [4].

²The code for the project can be found at <https://github.com/ufmg-smite/lean-smt>

1.3 Related Work

1.3.1 Hammering Towards QED

As previously mentioned, Hammering Towards QED is a project that aims to describe all the tools, which the paper calls “hammers”, that were created with the purpose of connecting automatic and interactive theorem provers. Besides that, this document also outlines the main components that such tools usually have. They are the following:

- The premiss selection module: that identifies a subset of the facts previously demonstrated in the ITP that are more likely to be useful in order to prove the given goal, to be dispatched to the ATP.
- The translation module: that builds a problem in the language of the ATP that corresponds to the original goal from the ITP and using the premisses that were selected.
- The proof reconstruction module: that lifts the proof produced by the ATP into a proof that is accepted by the ITP.

In our case, we will restrict ourselves to implement a proof reconstruction module. The three main strategies used to reconstruct the proof produced by the automatic system inside the interactive one are also described in the paper.

The first one is to use the ITP to verify a deeply embedded version of the proof received, and, in case it is succesful, reflect this proof inside it’s checker, proving the original goal. More specifically, the hammer defines a datatype to represent terms in the ATP and a set of functions to manipulate values on those datatypes, representing the axioms that the solver uses to transform the terms. Then, a lifting function is defined, that is, a function that take a value of this datatype and outputs an equivalent term in the native language of the ITP. Finally, the correctness of each transformation function is verified with respect to the lifting function, in the sense that, if the input term was lifted to a value that is provable in the ITP’s logic, then the output term will also be provable. The ATP’s proof will be represented as a sequence of those transformations, and their correctness are proved *a priori*. When checking a specific proof, the only step that the ITP must perform is to compute the result of the application of all the functions in the solver’s output, and to check if the final term matches with the expected one. This technique is known as the Certified approach [8].

The second one is to match each axiom in the ATP's logic into a proved lemma or a tactic `[is it okay to use tactic here? should I introduce the term first? -tom]` defined in the ITP that works directly with native terms of the system. The proof produced by the automatic solver is then parsed into a sequence of applications of those lemmas and tactics and replayed inside the ITP. In this case, the proof is built on the fly and doesn't have its correctness guaranteed (it can fail in the middle of the process in case the ATP or the hammer did something wrong). On the other hand, this technique skips computations done over embedded terms, which have to be done by the Certified approach, having the potential to have a better performance. This approach is known as the Certifying approach [8].

The third one is to compile the proof into the ITP's source code. This implies generating an actual script in the native language of the interactive system that corresponds to the proof received. After the script is generated, it is possible to postprocess `[maybe copy the references that hammering towards use to talk about this? -tom]` it in order to make the proof easier to be checked, in a way that the final script can possibly ignore a large portion of the original proof. This approach can be inconvenient for very large proofs, as it requires that the script is stored in some filesystem. However, it has the advantage over the two previous methods of only requiring access to the ATP on the first time that the proof is checked.

In this project we will be using the second approach. We give more details about this decision in the later chapters.

1.3.2 SMTCoq

One notable example of hammer is SMTCoq [12]. It is a plugin for the proof assistant Coq [5] that can be used as a tactic to prove theorems via their encoding into SMT and by lifting proofs produced by the SMT solvers `veriT` [9] and `CVC4` [3]. The tool can support multiple solvers due to a preprocessor written in OCaml, that is able to turn the different proof formats they emit into an unified certificate in the Coq language, that will be used as input for the plugin. Our tool explores one specific format of `cvc5`, therefore, for now, it only supports this solver and is less flexible in this aspect than SMTCoq.

The Coq hammer follows directly the Certified transformations approach, described in the previous section. It has a set of certified functions representing the transformations that can be done in the theory of Equality and Uninterpreted Functions and Linear Integer Arithmetic, as well as resolution chains. The way that SMT solvers prove that

a proposition is true is by showing that it's negation is unsatisfiable [it's okay to use unsatisfiable without defining it first? -tom]. This kind of proof is parsed in SMTCoq into a sequence of applications of certified functions, which have to transform the negation of the original goal into a term that will be lifted to the *False* proposition in Coq. Once the ITP verifies that the sequence of steps produced by the solver indeed produces *False*, the original theorem can be automatically closed, as we described in the previous section.

SMTCoq also differs from our work as it implements a translation module, but it doesn't implement a premiss selection module.

1.3.3 Sledgehammer

The ITP Isabelle/HOL [14] has a similar tool, namely, Sledgehammer [7]. This system have multiple strategies for building a proof for a theorem using ATPs. One of them is to invoke several ATPs in parallel with the given goal as their query and parse their output into a sequence of applications of Isabelle's predefined lemmas and tactics that could, in thesis, prove the original goal. This falls into the second category of strategies for proof reconstruction presented in Section 1.3.1.

Another technique used by Sledgehammer is to use the proof found by the ATP's only to scan which lemmas were used, and then discard it and build a whole new proof, feeding the lemmas found into existing plugins for Isabelle that already perform automatic theorem proving for the ITP. In this case, Sledgehammer works as a premiss selector (described in Section 1.3.1) for the other plugins.

1.4 Organization of this document

Chapter 2

Formal Preliminaries

2.1 Satisfiability Modulo Theories

2.2 Lean's Type Theory

take a look at chapter 2 of smtcoq

- Falar sobre porque eh facil confiar no proof assistant
- Explicar que taticas extendem a linguagem mas nao aumentam o trusted core

2.3 Lean's Framework for Metaprogramming

Chapter 3

Certifying Reconstruction of SMT Proofs in Lean

3.1 Certified vs Certifying

3.2 Classical vs Intuitionist (?)

3.3 Tactics

3.4 The Complete Architecture

3.5 Skipping the Parser

Chapter 4

Evaluation

Chapter 5

Future Work

Bibliography

- [1] Haniel Barbosa et al. *cvc5: A versatile and industrial-strength smt solver*. In Dana Fisman and Grigore Rosu, editors, *Tools and Algorithms for the Construction and Analysis of Systems*, pages 415–442, Cham, 2022. Springer International Publishing.
- [2] Haniel Barbosa, Andrew Reynolds, Gereon Kremer, Hanna Lachnitt, Aina Niemetz, Andres Notzli, Alex Ozdemir, Mathias Preiner, Arjun Viswanathan, Scott Viteri, Yoni Zohar, Cesare Tinelli, and Clark Barrett. Flexible proof production in an industrial-strength smt solver. In Jasmin Blanchette, Laura Kovacs, and Dirk Pattinson, editors, *International Joint Conference on Automated Reasoning (IJCAR)*, Lecture Notes in Computer Science. Springer, 2022.
- [3] Clark Barrett, Christopher L. Conway, Morgan Deters, Liana Hadarean, Dejan Jovanović, Tim King, Andrew Reynolds, and Cesare Tinelli. CVC4. In Ganesh Gopalakrishnan and Shaz Qadeer, editors, *Computer Aided Verification (CAV)*, pages 171–177. Springer, 2011.
- [4] Clark Barrett, Pascal Fontaine, and Cesare Tinelli. The SMT-LIB Standard: Version 2.6. Technical report, Department of Computer Science, The University of Iowa, 2017. Available at www.SMT-LIB.org.
- [5] Yves Bertot and Pierre Castéran. *Interactive Theorem Proving and Program Development - Coq’Art: The Calculus of Inductive Constructions*. Texts in Theoretical Computer Science. An EATCS Series. Springer, 2004.
- [6] Jasmin C. Blanchette, Cezary Kaliszyk, Lawrence C. Paulson, and Josef Urban. Hammering towards qed. *Journal of Formalized Reasoning*, 9(1):101–148, Jan. 2016.
- [7] Jasmin Christian Blanchette, Sascha Böhme, and Lawrence C. Paulson. Extending sledgehammer with smt solvers. In Nikolaj Bjørner and Viorica Sofronie-Stokkermans, editors, *Automated Deduction – CADE-23*, pages 116–130, Berlin, Heidelberg, 2011. Springer Berlin Heidelberg.
- [8] Valentin Blot, Louise Dubois de Prisque, Chantal Keller, and Pierre Vial. General automation in coq through modular transformations. *Electronic Proceedings in Theoretical Computer Science*, 336:24–39, jul 2021.
- [9] Thomas Bouton, Diego Caminha B. de Oliveira, David Déharbe, and Pascal Fontaine. veriT: An Open, Trustable and Efficient SMT-Solver. In Renate A. Schmidt, editor,

- Proc. Conference on Automated Deduction (CADE)*, volume 5663 of *Lecture Notes in Computer Science*, pages 151–156. Springer, 2009.
- [10] Davide Castelvechi. Mathematicians welcome computer-assisted proof in “grand unification” theory. *Nature*, 595, 06 2021.
- [11] Leonardo de Moura, Soonho Kong, Jeremy Avigad, Floris van Doorn, and Jakob von Raumer. The lean theorem prover (system description). In Amy P. Felty and Aart Middeldorp, editors, *Automated Deduction - CADE-25*, pages 378–388, Cham, 2015. Springer International Publishing.
- [12] Burak Ekici, Alain Mebsout, Cesare Tinelli, Chantal Keller, Guy Katz, Andrew Reynolds, and Clark Barrett. Smtcoq: A plug-in for integrating smt solvers into coq. In Rupak Majumdar and Viktor Kunčák, editors, *Computer Aided Verification*, pages 126–133, Cham, 2017. Springer International Publishing.
- [13] The mathlib Community. The lean mathematical library. In *Proceedings of the 9th ACM SIGPLAN International Conference on Certified Programs and Proofs, CPP 2020*, pages 367–381, New York, NY, USA, 2020. Association for Computing Machinery.
- [14] Tobias Nipkow, Lawrence C. Paulson, and Markus Wenzel. *Isabelle/HOL: A Proof Assistant for Higher-Order Logic*, volume 2283 of *LNCS*. Springer, 2002.